UBS GLobal Technology Conference

Company Participants

- Alysa Taylor, Corporate Vice President of Azure and Industry Marketing

Other Participants

- Unidentified Participant, Analyst, UBS

Presentation

Unidentified Participant

Okay. Thank you, everybody. I'm honored to lead this keynote with Alysa Taylor of Microsoft. I wanted to just thank the whole Microsoft Team. You've been fantastic partners at the event this year. I don't think it would surprise anybody that Microsoft and NVIDIA are the two most requested companies for one-on-one meetings. I think that says a lot about where we are in the tech curve. So, thank you Brett, Kendra, Mary, and the team for making yourself available for so many one-on-ones. It's fabulous and elicit for making the trip to Scottsdale to chat AI and a number of other topics with us.

Alysa Taylor

It's a pleasure to be here. Thank you.

Unidentified Participant

Oh, great. Good. Alysa, do you want to take a minute and share with the group your role at Microsoft?

Alysa Taylor

Absolutely. So, I'm responsible for our Azure business. So, that comprises all of our AI services as well as our data infrastructure, digital app, and innovation and infrastructure. So, that's the entirety of the Azure portfolio. And then I also have the responsibility for our global industry. So, everything from our regulated to non-regulated industries including sustainability.

Unidentified Participant

Okay.

Alysa Taylor

And that's all in the go-to-market side. So, bringing all of the things that we build via partner, how do we then construct that in a way to bring them to market.

Unidentified Participant

If this was a two-hour keynote, we could fill it up with all of your expertise. We'll try to condense it to 28 minutes. So, let's hit on a couple of hot topics to start.

Alysa Taylor

Okay.

Unidentified Participant

So, I think this might be Microsoft's first occasion to at least sit down with the investor community more formally post the drama of a couple of weeks ago with OpenAI. So, understanding that you are probably somewhat limited in what you can say, I'd love if you could share some perspectives. I can share with you that some of the questions we received had to do with, you know, IP ownership, it had to do with the extent to which Microsoft might be better off, assuming they are, given how things settled down. But I'd love your perspective, so if don't mind (multiple speakers)

Alysa Taylor

There was drama with OpenAI. But yes, no, it's been -- it's been an interesting week. You know, what I would say is, we are, you know, committed to our partnership with OpenAI. We are excited to see that Sam has been reinstated as the CEO. And we have a great roadmap of innovation ahead and we have perpetual use rights for the IP with OpenAI and we have everything we need to fulfill on our roadmap of innovation.

Unidentified Participant

Okay.

Alysa Taylor

So, business business as usual.

Unidentified Participant

Okay. Nice to hear. Let's talk about another big event recently and that is, I think you played a central role with Guthrie and the team at Ignite, so that was a big event. We had one of our team members there to talk with customers and partners and learn. But what were the couple of key takes for you coming out of Ignite?

Alysa Taylor

It was a really exciting event for us, because we introduced a bunch of innovation and I would say across the entire stack of the Microsoft cloud. And so, some of the highlights that I would bring forward, our big announcement around silicon I think was a pretty prevalent part of the event was our first entry into first-party silicon. We introduced Azure Maya, which is our AI accelerator.

We also introduced Azure Cobalt, which is our cloud-native CPU. But that really is part of our larger, what we call silicon to systems. So, we also had -- we announced the general availability of Azure Boost, which is the technology that allows you to offload networking and security onto purpose-built hardware, so that improves the performance. Also, our hollow fiber core which improves networking. So, lots of infrastructure.

And then, we also announced the general availability of Microsoft Fabric, which is our new analytical data platform. This is a very important key asset in our AI services because it allows organizations to be able to take a bunch of -- a lot of disparate data from either on-prem, different data sources, different cloud providers, Snowflake, GCP, to be able to bring that data into an integrated data environment called OneLink. And then we have native AI services that are built into Fabric. So, really exciting announcement there.

And then when we think about the tooling aspects, so we also brought the Azure AI Studio into preview. And that includes all of our responsible AI tooling like content safety, our Model-as-a-Service, so that you can actually govern and reason over third-party models. And then, Azure Arc which is our central control plane that allows an organization to do the management and governance of all of their cloud and on-prem assets. So, kind of this next-generation of hybrid. We introduced what we call the adaptive cloud which allows organizations to manage in a central integrated control plane all of their on-prem multi-cloud cloud assets and to be able to project those assets to their edge environment. So, it's bringing together all of the decade innovation that we've had around Edge, IoT, as well as application management into one single source.

Unidentified Participant

Okay. Let's, I mean, can we unpack the first one a little bit?

Alysa Taylor

Of course.

Unidentified Participant

You led with it, so let's go there, the new Maya chips. So, I think, one question on everybody's mind is, Microsoft has an enormous and successful partnership with

NVIDIA. So, I think everyone's question was, how does Maya fit into that relationship? Is it targeting a different workload type? Maybe you can elaborate a little bit on that.

Alysa Taylor

Yes, it's interesting. The AI workloads are, like HPC and AI workloads are different than enterprise workloads. Because you actually, when you have these workloads come in, they are there asynchronous. So, you can have them at -- both at peak times, burst capacity or not at all. And so when we talk about the silicon to systems approach, the systems approach is really allowing us across a network of hardware to be able to load-balance, so that we can optimize for peak performance as well as cost. And so for us, you know, our first-party silicon was just yet another addition to what we have built with AMD, and Intel, and NVIDIA. And so, that's the systems aspect of it. So, it allows us to be able to control the performance of those workloads across the variety of the systems and that's where things like Boost come into play, as well as the innovation that we've seen in the Fiber layer as well.

Unidentified Participant

So, there's a cost benefit to Microsoft as well internally.

Alysa Taylor

There. Yes, so if you think about it, you know, GPU resources are very resource capital intense. And so you don't want them to go unused. And so, this allows us, as you it is opaque to our customer, but allows us to be able to balance the workload on the backend across. So, we have no plans to replace our NVIDIA or AMD investments. It's just adding to the ability to do the load balancing.

Unidentified Participant

I'm assuming you spent Ignite talking about AI endlessly (multiple speakers)

Alysa Taylor

Yes.

Unidentified Participant

With your customers. So, I think everybody could also benefit from your perspective on how those conversations are going, you know, where we are on the adoption curve, where customers that are in the use-case discovery process. What are those use cases that are coming up in that initial wave? In other words, where are they getting the greatest ROI near-term? That might be helpful for everybody.

Alysa Taylor

The use cases that we see emerging, really, we look at it in three dimensions. So we look at it by very specific industry use cases, what are the line-of-business use cases, and then what are the individual function use cases? And so I'll give you just a couple of examples of some of the more prevalent, that we see. In industry, healthcare is one where one of the greatest challenges that healthcare organizations face is physician burnout. So, we actually have seen in the US, physician burnout has increased from 43% to 52% in the last three years. So, that's the post-pandemic. We're still seeing physician burnout on the rise.

And so with technology like DAX Copilot, it is a combination of ambient and generative AI that can record an interaction between a physician and a patient, whether that be in the telehealth or an exam room visit, it then automatically analyzes, summarizes, and generates a clinical note that can be uploaded into the electronic medical health record on behalf of the physician. So, that reduces the administrative burden that a physician has. And you see organizations like Atrium Health that have deployed DAX Copilot. They have actually recorded that their physicians are on average getting 40 minutes a day back, and they're improving the overall care of their patients. And so that's a great use case of using AI technology to combat a very critical industry problem.

In the horizontal space, customer service is one that we are using, we're seeing that as one of the most prevalent use cases. Internally at Microsoft, we have deployed the Microsoft Copilot for service. And in a matter of months, we've been able to increase the first call responded by 31%, and actually improve satisfaction by 12%. So, you see these use cases where the technology is being deployed and there's real outcome associated with it.

And then in the function, a great example of this is cyber professionals. And if you think about the sheer amount of data that a cyber professional has to reason over to be able to detect threats, Microsoft Copilot for security actually aids in that process of being able to come through all of the potential threats, pinpoint the most high targeted threat and then allow about cyber professional to actually work against that -- that threat. And so, that's where you see sort of industry line-of-business and function.

And, you know, in addition to the use case, I often get asked, what is the return on this, what is the actual tangible monetary return, and we worked with IDC. We surveyed 2000 global companies. And there were some really interesting stats that came out of that. The first is that, over 70% of the organizations had already deployed AI technology. So, that sort of shows the interest level of AI technologies. 91% of those organizations had indicated that they had an AI project up and running in less than a year.

So, if you think about just we've talked for years about digital transformation, that speed. So, it's just the pervasiveness of the adoption, the speed of adoption. And then probably most interesting to this community is that for every dollar -- it was reported for every dollar that the organization had spent on average, they were seeing a return of $3.5 return on that dollar.

Unidentified Participant

Wow.

Alysa Taylor

So, that's -- you know, that's a really -- so it's taking, those use cases are now starting to see clear patterns of (multiple speakers)

Unidentified Participant

Yes.

Alysa Taylor

ROI on them.

Unidentified Participant

Interesting. Let's talk a little bit about how Microsoft is scaling its AI infrastructure to meet that (multiple speakers)

Alysa Taylor

Yes.

Unidentified Participant

Incredible demand. So, I think there's a perception that there is a significant supply shortage of NVIDIA chips out there. Can you comment broadly on the severity of that supply constraints and how Microsoft, you, Satya, the leadership are scaling up that infrastructure to meet the demand?

Alysa Taylor

Absolutely. When we think about capacity planning, it's very multifaceted. So, we start with demand planning, so that is both the short-term demand and then projections for long-term demand. And then we couple that with resource allocation. And so, resource allocation is both the availability of the infrastructure, but then we also take into account things like geo-availability, LAN, power. So you can see it's a very multifaceted approach to how we think about capacity planning.

And really the systems approach that I keep talking about is so core to this because that is how we are able to be able to do capacity planning across our GPU allotments, whether that be through our partnerships across NVIDIA and others, as well as now our own first-party investments in it and then being able to do demand shaping by geo. And, you know, we do this all wrappered in our commitments to

sustainability. So, we will actually come -- we will need to be on path for our 2025 carbon-negative footprint within our data centers. And then we also have a goal by 2025 to protect more LAN than we use in our data centers.

Unidentified Participant

Okay. And Alysa, as Microsoft works to stand-up that AI infrastructure as fast as you can.

Alysa Taylor

Yes.

Unidentified Participant

I recall a blog post in August. I think Satya and the team were talking about getting live on a number of H100 clusters and that there's hundreds of thousands of H100s that are coming live over the next year.

Alysa Taylor

Right.

Unidentified Participant

So that's a big revenue unlock, it feels to me, for Azure. So, one thing I'd love to understand a little bit better are what the constraints are to get that GPU supply ready. You mentioned a few of them.

Alysa Taylor

Yes.

Unidentified Participant

Not only getting your hands on the chips, but the power requirements are different.

Alysa Taylor

Absolutely. Right.

Unidentified Participant

The -- literally, the LAN, the data center, the networking architecture, are any one of those constraints any greater than the other in terms of standing up that infrastructure, or is it a collection of all of those?

Alysa Taylor

It's a collection of all of it. I mean, I think that's why we take that integrated approach to capacity planning. You have to, because it is both the infrastructure, the geoavailability, the land use, the power structure, it's all of it. And you mirror that against demands today, and then the forecast against demands, and then you demand shape, like I said, you then sort of help. You know, we have different data centers where we will -- certain workloads are more, you know, attuned to a certain geo.

There are more sustainable data centers than others. And so, we work directly with customers to make sure that they are in the best data center possible. And then we, on a daily basis, monitor capacity of those both data centers and we have a weekly socket that we review across our leadership team on things like GPU availability and capacity. So, it's a very hands-on integrated approach.

Unidentified Participant

Without giving up any of your secret sauce, AWS and Google are also trying to stand-up their AI infrastructures as fast as possible. Are there any differences in the way that Azure is constructing its new AI infrastructure that might give it a competitive edge a couple of years out?

Alysa Taylor

You know, I think the thing that we point to, we'll continue to point to, is the entirety of the system. It is not just about the chips. I think people like to talk about the chips. We have for many years had deep partnerships across Intel, AMD, and NVIDIA. We continue those partnerships. You probably saw yesterday, AWS just entered into a partnership with NVIDIA. That is -- we have a long history with NVIDIA, and it has been part of how we've built out our infrastructure. But it's every aspect of it. It is the networking, it's the security. And we're building all of this to be able to be the most performant and the most cost-efficient for customers.

And we actually -- there is a benchmark of supercomputing in the industry of 500 cloud providers and we are number one in this supercompute category. And so we feel like this integrated approach to both how we build out the infrastructure as well as how we do capacity and demand planning, is a real benefit and it's paying off not only in industry benchmarking but also for our customers to have the most performant and cost-effective. Okay, let's talk about the pricing a little bit on those Azure compute resources. Yes.

Unidentified Participant

A lot of us watched Sam Altman and the team at the recent OpenAI Developer Day.

Alysa Taylor

Yes.

Unidentified Participant

There are a number of cool announcements from that day. But one thing that struck me as really interesting is that Sam talked about essentially reducing the per token price of their premium models.

Alysa Taylor

Correct.

Unidentified Participant

Like GPT-4.

Alysa Taylor

Yes.

Unidentified Participant

Like by an extraordinary amount.

Alysa Taylor

Correct.

Unidentified Participant

3x.

Alysa Taylor

Yes.

Unidentified Participant

So, my first reaction was, boy, if they could reduce the price per token by 3x, knowing that the consumption of Azure resources are a big part of that cost structure, does that mean that LLMs are becoming incrementally less compute-heavy, compute-centric? Which probably wouldn't be a great thing for Azure. How would you respond to that?

Alysa Taylor

Well, I'll start by saying, we are obviously -- we continue to be very bullish on our AI growth. But the misnomer in the industry is that AI is an enterprise workload. And the reality is that digital natives, small businesses, benefit from the power of AI services. And so by being -- by making the models more efficient, we are -- and then bringing down the price, we are actually opening up our addressable market.

Unidentified Participant

You're almost like democratizing it.

Alysa Taylor

Correct.

Unidentified Participant

And it's a P times Q phenomenon.

Alysa Taylor

P times Q. Correct. And so it allows us to be able -- that efficiency allows us to provide these services to more and more customers where the barrier to entry before was price. So, efficiency in AI services is actually a good thing because, exactly to your point, it allows us to democratize the services to more and more users.

Unidentified Participant

Maybe a similar question. Your answer might be the same, but another question that investors have is that, as the bulk of compute needs changes from training of the models to inference of the models, is that less compute heavy on the inference side, such that as we go through that training to inference shift, does that change the resource needs for Azure?

Alysa Taylor

Well, there's this new emerging category, right, which is, you know, LLM Ops. It's, the ops of it. And so, how do you do things like prompt engineering, rag, fine-tuning, all of that, we actually -- one of the announcements that we made at Ignite was Model-as-a-Service. So, to really bring -- so that isn't a very labor-intensive, we're actually wanting to bring more and more tooling to the LLM Ops space, which is a good thing, because then it allows people to train more models, use more models, have kind of the different models for different use cases and be able to reason and govern over those models.

So, there's the compute aspect of the models and the inferencing of the models, but then there's the data aspect of it. And so more and more, you know, your AI is only as good as your data estate. And so, you see it's not just about the compute of the

models, but it's also the ancillary services that come with when you are building or modernizing an app, data being a key component of that, storage being another key component. And so, you know, there is a -- we look at it from the AI platform side, compute, data services, networking and storage.

Unidentified Participant

I'm going to ask you about the data services in a quick (multiple speakers)

Alysa Taylor

Oh, good. This is one of my favorite topics.

Unidentified Participant

First, let's summarize this up to the Azure level that Bret and his team conveyed to the street. So, Amy on the last call gave guidance for stable Azure growth after the December quarter. And there's a couple of dynamics in there. It feels as if it's reasonable to conclude that the AI contribution to Azure's growth should improve. And so, stable means that the core or ex-Azure part of -- ex-AI part of Azure might decelerate slightly, maybe that's scale economies. But can you comment a little bit on that guidance? Why would ex-AI Azure growth moderate? Really, is it just -- is it scale?

Alysa Taylor

Well, if you think about the sheer number of workloads that you have to bring online to be able to deliver that stable growth, which is, you know, our forecast for H2. And then the other thing that I would say is, you think about Azure from the entirety of the Azure platform, right? You have the infrastructure layer that we've talked about. You have the management and governance layer, you have the data layer, you have the -- all of the -- you know, the AI services, the tooling, everything that we provide at the app layer. So, there is in the workloads that make up Azure are fairly, you know, in comprehensive and significant.

And so, when we look at Azure growth, it is both around growing new workloads and then you'll see customers do, which is normal course of business, which is continuing to optimize those workloads for performance and cost. And so, you know, our commitment is to grow workloads across every aspect of the Azure Cloud while making sure that we are working hand in hand with our customers to, you know, maximize their price to performance.

Unidentified Participant

Okay. Can we hit a little bit before we get to data on the optimization side? I think there's acute interest in that.

Alysa Taylor

Absolutely.

Unidentified Participant

So, I'm sure at Ignite and through the course of your day, you're talking a lot with customers about that, amongst other things. So, where would you characterize Microsoft's Azure customer base in terms of how far along they are in terms of that journey to optimize their spend?

Alysa Taylor

Well, optimization is an ongoing occurrence because if you think about when you bring a new workload online, whether that be migrating a on-prem workload to the cloud, whether it be building a new application, you're constantly optimizing the performance of that workload. And so, optimization is just, it's an ongoing part of cloud infrastructure and cloud compute. And so, for us, there is two things that we want to do. One is, work very closely. We provide a number of resources, whether that be Azure landing zones. We have a service called Azure Advisor that actually helps customers build well-architected workloads and then how they construct those workloads in the most performant way. And that's part of optimization.

And so that, I think, is a normal course of business. And then as they -- you know, as they optimize, they bring on new workloads. And so our goal is to help customers build and construct the best workloads at the best performance and then make sure that we are working with them on that next workload.

Unidentified Participant

Okay.

Alysa Taylor

And so, that's where we, when we talk about workload growth, it is in conjunction with making sure that, you know, our customers are building well-architected services at the beginning.

Unidentified Participant

Okay. Got it. Let's talk about data now.

Alysa Taylor

Okay.

Unidentified Participant

A number of interesting topics here. So, one that intrigues me a little bit is this notion that in advance of, or perhaps concurrent with enterprises moving forward on AI

projects, they need, to, quote, get their data estates in order.

Alysa Taylor

Yes.

Unidentified Participant

Are you hearing that? And to what extent is it starting to pull through, even if it's just in conversations today, the rest of Microsoft's Data Suite?

Alysa Taylor

Well, the -- as I mentioned, your AI experiences are only as good as the data that the AI reasons over. And so, it is very much, you know, step one is what is the data set that you want to be able to apply AI services against? And that data set has to be in the cloud. So, there is kind of step one, which is being able to access or migrate your on-prem data and then organize it and construct it in a way that the AI services can pull from it. And so there is a number of things that we are doing to aid customers in being able to get their data estate in order.

One is, I talked about fabric. Fabric is for us a very meaningful investment in our AI pursuits. So, we -- Fabric is our analytical data service as I mentioned. One of the unique capabilities is the OneLake aspect of Fabric which allows you to bring in various different disparate data sources in a -- through shortcuts. So, you can shortcut into data, you can mirror competitive data, and we're competitive. So, you know, in different -- whether you have data in GCP or AWS, you can bring that into OneLake and then you have an aggregate data. And then we have applied our AI services directly within Fabric.

So, if you think about how you then -- you take that data and you organize the data to be able to call against very specific data sets, things like Vector search become very important. So, we've brought Vector search into Fabric through -- it's called AI search. We've also brought it into our distributed data scale, Cosmos DB. You can have AI search natively integrated into that. So, it's -- you know, it is both the aggregation of the data and then actually how you construct and access the data. And those are two really critical, important things. And so Fabric and AI search are very major investments that we've made as organizations are bringing their data into the cloud, organizing their data to apply the AI services on top.

Unidentified Participant

And also, I'd be curious, are you seeing any evidence that AI is starting to accelerate the pace at which organizations are embracing the cloud? Are they -- is it boosting the on-prem to cloud data migration effort? Is it causing any acceleration in the pace at which you're running your databases in the cloud versus on-prem? Is that beginning to happen? I'm sure it's been a constant process. I'm curious if AI is acting in any way as an accelerant.

Alysa Taylor

Absolutely. If you think about there -- when you think about AI in the terms of applications, you are either modernizing an application or you are building a net new application. So, if you even take Microsoft, what we've done with M365 Dynamics, GitHub, that's modernizing an existing application. And so, when you modernize that existing application, you're bringing in the Azure OpenAI service, which are the API calls into your data set. That's why your data is so important. But then you also see an increase in things like storage, because you actually have to store those processes. So, you see, when it is not just about the API calls, but it's about the aggregation of that data, the data services increase, as well as ancillary services like storage.

When you get into build, that is about kind of the entirety of the stack that you build upon. And that is everything from, you know, how you -- the developer services that you use. So, whether that be GitHub and the GitHub repos, GitHub Copilot to actually do the coding, then bringing data into a operational data service, so Cosmos DB, as I mentioned, is our distributed at scale database, then being able to apply the app services on top of it, as well as the developer ops management. And so you have this, you know, to create a net new application, you have an entirety of services that you construct to build a net new modern AI application. And so, both in the modernize and build, we're seeing really unique pull-through across the Azure platform.

And so, a great example is UiPath is a customer. They modernized their business process software with Azure OpenAI. Their storage and networking increased. They then started to build out net new applications. So, building their data services as they brought new data services in to build out new applications increased.

And then another great example is the customer, Real Madrid. So, they are continually trying to figure out how to engage with their fan base. They were able to originally only call from seven data sources. Now they are able to actually -- or sorry, five data sources. They're now able to call 70 data sources by being able to do that data aggregation that we talked about in OneLake. And then, a really unique thing is they've actually increased their fan profiles by 400% in two years. So, then it's a nice marriage of like, our ACR increases and most importantly, our customer business outcomes increase as well.

Unidentified Participant

This might be a good time to ask you a little bit about the Oracle relationship on the database on Azure. So, by the way, Alysa, your profile among investors went massively up when you were seated in between Satya and Larry for that session. I think everybody in the room probably watched you then. How fun must have that been, right?

Alysa Taylor

It was very -- it was an incredible experience.

Unidentified Participant

Rivalry and defense [ph]

Alysa Taylor

Nerve-racking to have the two of them watch you as you deliver their intro. But it was an incredible experience. It was Larry's first time in Redmond.

Unidentified Participant

Yes.

Alysa Taylor

Which was really kind of a momentous experience.

Unidentified Participant

Yes. And it's a big deal, not just for investors. But for large organizations, UBS. It's public information. We're an extraordinary large Azure customer. But we also have a massive Oracle database.

Alysa Taylor

As most customers do.

Unidentified Participant

Yes, exactly. So, it mattered a lot to us. So, the question to you is, what does that mean for Microsoft, how are you going to benefit from this partnership on the database side?

Alysa Taylor

Well, as we've been talking about the -- how data -- how critical data is to building out unique AI experiences, the partnership that we have with Oracle was actually a direct request from our customers of being able to bring their Oracle databases into Azure. And so, we have a very unique partnership where you can run your OCI database directly in Azure, so as you're modernizing and that governance management and modernization can happen in a very integrated way. And for us, you know, it gives, you know, Oracle's over 400,000 -- over 400,000 customers, the ability to bring OCI directly into Azure and run natively in an Azure environment.

And so, that's a very -- you know, we've worked to be able to architect that in a way that is most beneficial to customers. So, they can just literally take OCI, run it in Azure, and then have all of the ancillary Azure services that we talked about, to be able to do the management, governance, and modernization.

Unidentified Participant

Yes. Huge partnership.

Alysa Taylor

Huge partnership. And the incredible thing is, we're seeing a lot of interest from organizations come forward for this reason to say, like, it has been prohibitive for them to be able to bring their Oracle databases to the cloud. And now this allows, you know, that way for them to very seamlessly bring it into an Azure environment.

Unidentified Participant

Let's talk about a couple of the applications that bring a lot of your AI capabilities to the customers. So, the Copilots.

Alysa Taylor

Yes.

Unidentified Participant

So, we'll ask you about two. So, on GitHub Copilot.

Alysa Taylor

Yes.

Unidentified Participant

Color Alyssa on how the traction is progressing, what kind of developer productivity gains you're seeing in these early days?

Alysa Taylor

So, we talked about kind of the use cases by industry, by line of business by function. Developer is one of the key functions, that where we see enormous potential. There is a -- and this has been happening for years, but it's what's called the app gap in the world which is, there are more applications being wanting to be built than there are developers. And so, you -- this is why you see things like the low code tooling come online. And so, it is both about making the developers more productive and giving them the tools to be able to do that. And so GitHub Copilot has actually been a phenomenal productivity gain for developers. So, developers that are using GitHub

Copilot have a increase of 55% productivity. That's an enormous number if you think about like making developers 55% more productive. And we have, you know, and that base of users is growing. So we have over a million GitHub Copilot users, paid users. So, it is a tool that has become very instrumental to developers and the -- you know, there is the making them more productive, but are they happy about it. And you actually see this really nice marriage of, they are more productive and they have reported an increase in satisfaction in their day-to-day work as well.

Unidentified Participant

Okay.

Alysa Taylor

So, that for us is a sign of true success, productivity plus gains and satisfaction.

Unidentified Participant

Got it. Let's now talk a little bit about M365 Copilot. Only been GA for a couple of weeks, so, understanding it's super early.

Alysa Taylor

Yes.

Unidentified Participant

But what can you share with everybody listening in and here in the audience about the traction that Microsoft has seen so far? Can you share anecdotes or any?

Alysa Taylor

There's a couple of things.

Unidentified Participant

Probably not metrics, but anecdotes.

Alysa Taylor

Yes. These guys will not -- you know, will not let me share too many metrics. But the -- as you mentioned recently, GA Copilot intentional from a naming standpoint is that it is a Copilot for individuals, right? So when we talk about GitHub, when we talk about M365, it is about enhancing the productivity and both enhancing the productivity and what we call unleashing the creativity. So allowing humans to be more creative and more satisfied in their work. We at Ignite released our Work Trends Index which reported of the Microsoft 365 Copilot users, 70% of them, they indicated that they were more -- that they would not give up Copilot, that it had become instrumental

into their day-to-day environment, and that it had made them more satisfied in their work. And so, you see this, you know, for something that is so early to market, to see that large of a base say that it is critical to their day-to-day work is pretty impressive.

Unidentified Participant

The UBS CTO has promised that I'll have one eventually. I'm waiting. So, as soon as it can arrive, I'll welcome it.

Alysa Taylor

And I will tell you, PowerPoint summarization.

Unidentified Participant

Yes.

Alysa Taylor

Is key.

Unidentified Participant

Okay?

Alysa Taylor

And I highly recommend that. If you -- if any of you get PowerPoints like I do, which is dense and large volumes of PowerPoint slides, there is a summarization button that can take those very dense large number of slides and summarize them into a pithy summary for you.

Unidentified Participant

I look forward to that.

Alysa Taylor

Quite key.

Unidentified Participant

Let's give you a chance to make some closing remarks. I know we're out of time. So, maybe, Alysa, but when you think about the next three-plus years managing a business as broad as Azure, what are the couple of things that gets you most excited in terms of the growth trajectory?

Alysa Taylor

Well, I would say, the moment in time that we're in depending on the age of some of us, like we've lived through some pretty big shifts in the technology and if you think about internet, mobile cloud, we're now with this next inflection point which is this AI transformation and you've seen adoption happen at a rate that I have never seen before in terms of just the ground swell of people realizing and evaluating the potential of what AI can do both for their company, for their individuals, and I just think it's we're in a really unique point in time in the market and so it's exciting to be here and so it's not only just as we've talked about the individual AI services, but it's the ability across different industries, different functions, different business units to really redefine how they work and how organizations interact with our customers how they reinvent business processes. So, for me, the most exciting thing is just this inflection point that we are at and what -- and we are early, early days, very early days of what's possible.

Unidentified Participant

You're going to have an exciting 2024 I predict.

Alysa Taylor

Think so.

Unidentified Participant

Microsoft, Alysa. Thank you so much for attending the event. I enjoyed that conversation.

Alysa Taylor

Thank you very much.