

Evercore ISI Inaugural TMT Conference

Company Participants

- Jensen Huang, Founder, President and Chief Executive Officer
- Manuvir Das, Head of Enterprise Computing
- Simona Jankowski, Director of Investor Relations

Other Participants

- C.J. Muse, Evercore ISI

Presentation

C.J. Muse

Well, good morning, good afternoon, everyone. Thank you for joining us. My name is C.J. Muse -- and Research for Semiconductors and Semiconductor equipment. Thank you for joining Evercore ISI's Inaugural TMT Conference, this afternoon. Quite pleased to have Nvidia and more specifically Manuvir Das, Head of Enterprise Computing, which -- he's s quite timely given what we've heard most recently at Computex as well as I guess, a few months back at GTC. We also have Simona Jankowski on the line Director of Investor Relations.

And with that, let me turn it over to Simona real quick.

Simona Jankowski {BIO 7131672 <GO>}

Thank you so much CJ for hosting us. Before I turn it over to Manuvir, I just wanted to very quickly cover our forward-looking statements. So as a reminder, this presentation contains forward-looking statements, and investors are advised to read our reports filed with the SEC for information related to the risks and uncertainties facing our business.

Okay, with that, I'll turn it over to Manuvir.

Manuvir Das

Yeah, thank you very much CJ and Simona. I thought what I would do CJ at the outset is just talk a little bit about the background for what we're doing and why we're doing it, right? You referenced the work we're doing with enterprise customers and what we talked about at Computex and I think it really starts with the question of Jensen our CEO has taught over the last few months about the coming democratization of AI and really what that means is, where is today with the work

we've done over the last decade, we have thousands of companies around the globe using artificial intelligence in Nvidia technology when we talk about democratization of AI we're talking about tens of thousands of thousands of companies all using AI, that's what we mean by democratization and I think the conversation really begins with why? Why is it that we believe that all of these companies are actually going to adopt AI, and I would like to share our point of view with this slide here. There's three reasons why we think every company will use AI. I'll focus on a couple of them.

The first is that, when you democratize a technology across a wide base of enterprise customers, really most of the customers who use the technology will not be directly aware of implementing that technology. So whereas today's customers are very aware of AI and they're putting in the infrastructure, the hardware and the software that it takes to do AI, they have their own data scientists, who develop things when AI is democratized, every company will be using it because of the applications that they use, which will be built with AI. A great example is Microsoft Office, which is infused with AI. The work Nvidia's doing with Cloudera, which is a data analytics platform that has now been infused with our technology and for the user, there is actually no difference in the experience. It's just that there's some special hardware underneath in every server, and the software has all been updated. So we see a big shift coming where every application will be infused with AI to make it more intelligent, but the customer's experience at the average customer company will be that it's the same application that used before. So that's one aspect.

The second aspect I want to highlight is on the right hand side of this picture, which is AI in every product and if you think about where we are at today, early adopters of AI would really internet companies, people with a large online presence, and they do things like for example, how do I recommend to you the next book to read, the next movie to watch, the next product to buy, I have a big social audience, I'm interacting with them but more and more now, we see companies that just built products regular products, where AI is infused into the life cycle of the product how do I design the product, how do I manufacture the product, how do I engage with my customers, how do I learn from that engagement and improve my product for the next cycle. So that's another democratization opportunity where every company regardless of what it does, would be utilizing in AI.

So this is the shift that we already see underway and really the work that we're doing on with the hardware and the software is to prime the entire ecosystem from top to bottom to participate in this democratization, so that every company can benefit. So with that, I have a couple slides just to lay out the work that we've been doing to enable this. This picture here represents the full platform, Nvidia is a full-stack computing company. Today we are much more internally a software company than we are a hardware company, we have many more engineers working on software than and hardware, and the reason is because when you use artificial intelligence, you use it in a variety of different use cases, whether it's video analysis, or conversational AI, or cyber security, or robotics, and you need a powerful software stack in each of those situations to enable that capability so an application developer can just complete the solution.

So if you look at the stack here, it starts at the bottom with servers, mainstream servers, volume servers, the same servers that are produced by the OEMs in the world today that are racked into enterprise data centers every day but now, with a little bit of GPU from Nvidia in every server. Then we layer on top of that, VMware vSphere, VMware as the de facto if you will Operating System of Enterprise Computing. It's got a very a very large footprint and then we've taken all of the essential software that we've built over the years, the frameworks, the toolkits, the SDKs, almost the operating system of AI if you will and we've packaged it together into something called NVIDIA AI Enterprise, which enterprise customers can use like any other platform that they deployed to-date and then all of the frameworks that we've built for these different AI use cases, and so with this, we have a complete computing stack from the hardware to the software, making it much more of a turnkey experience for the enterprise customer.

And then because of this shift now to this model, whereas we have been working very closely with OEMs on best-of-breed servers for AI -- bespoke servers, you do one thing with them, you do AI with them, they're relatively -- they have much higher price points because of all the hardware inside with this motion we're talking about here with the democratization we worked with the OEM's to really come into the mainstream servers as I said and so, if we go to the next slide, what we are very excited about now is, we've taken all our technology and we really rightsized it for the mainstream. We've made smaller GPUs more, power-efficient GPUs, more cost -- lower cost GPUs and CPUs with our BlueField and then we work with all of the OEMs to produce systems with these servers in them and we recently announced a whole wave of new systems from new vendors. I'm proud to say that, as of now we have 103 different server and workstation configurations from 16 different manufacturers and that number keeps going up.

So this is really CJ, the backdrop to the conversation, which is we've been sort of in a phase where of the last decade we've developed the hardware and software for AI. We've got the traction we -- our business has been doing well, we have thousands of customers using our technology, but now we're really shifting to the democratization where it becomes tens of thousands and hundreds of thousands of customers.

Questions And Answers

Q - C.J. Muse

(Question And Answer)

Well, that's a great start. Thank you for that. So I guess interestingly you talked about NVIDIA becoming more of a software company, particularly bringing AI to the masses. So I was going to say actually start there. Jen, you announced a number of partnerships or platforms whether it's VMware, Omniverse your newly announced base command AI software development platform. I think a lot of this is underappreciated by the market, so I guess let's take a little bit deeper here. So let's start with VMware with VMware partnership and your AI Enterprise software suite. How do you think about the percentage of sockets that would both benefit and be

willing to take advantage of this type of Enterprise license and how do you think about the TAM expansion for NVIDIA's Enterprise Hardware business?

A - Manuvir Das

Yeah, I think that's a great question to start CJ. So, the first point I would make is clearly, it's a TAM expansion for NVIDIA because this is new business we're entering with monetize software for Enterprise customers, right? We think the opportunity is very large because we think that any company that has an estate in their data center that their IT admins are managing a VMware environment and sharing our resources across applications and users, there'll be an opportunity for them to farm out AI resources both for the applications and for the in-house data scientists. So it's a for both the R&D side and the production side of AI. It's a very large opportunity if you just look at the simple math there's 8 million enterprise server CPUs a year and then we have a pricing model now for NVIDIA AI enterprise software its \$3,600 per socket on a perpetual license there's an additional 25% of annual maintenance. vSphere which this is attached to has over 300,000 customers. So it's billions of dollars of opportunity over time and we're really looking forward to opening up that market for the first time for NVIDIA.

Q - C.J. Muse

Really helpful and so I guess, could you also speak to the enablement and deployment of AI applications through the suite.

A - Manuvir Das

Yeah. I think the way we think about that CJ is really that, it's a three layer stack, when you deliver technology like AI. The first layer that people associate with NVIDIA, which is where opportunity tradition is been are on the hardware, right, the chips in the systems and then what I just talked about was the second layer, NVIDIA AI Enterprise, which is the operating system of AI, it's all the essential things you need to have a platform for AI, which I just described. But then the third opportunity, the third layer is all those applications on top, which NVIDIA has invested in very heavily over the years, whether it's CLARA for healthcare, Metropolis for video analysis of the Edge, and we certainly think that there is a significant opportunity there as well and as we go forward over time, you can -- you will see us opening up those opportunities and what's interesting about that level CJ is that for each of those applications, what drives the opportunity is, what is that application doing. So for example, is there a highly trained highly paid doctor that has to spend less time analyzing X-rays because this application software is helping them get started with it that's one kind of value, right? Am I processing something at a store where I'm being able to control the amount of materials that goods that are inadvertently taken out of the store that's a different kind of financial calculation for the customer, right? So, I think at the application layer, we do them one at a time, and it's really based on what is the value that is being provided to the customer in terms of the application.

Q - C.J. Muse

So just I guess -- a little bit deeper on the application layer. Is there a license opportunity there as well and I guess perhaps bigger picture, as you think about

your business entirely better price over time what percentage of revenues do you think will become recurring?

A - Manuvir Das

Yeah, I think there is absolutely a licensing opportunity and that's what I meant when I talked about the value, right and the licensing opportunity for each of these applications is different and it's dependent on what is the value that is being provided to the customer and how complicated is the problem that we are solving on their behalf, right? So we definitely anticipate that. I think as far as your question about the -- how we look at that opportunity compared to the hardware. I think, certainly over time, we expect that the software opportunity is at least as large as the hardware opportunity. The hardware opportunity already in the enterprise we believe is well over \$10 billion and we've talked about that before. We see the software opportunity certainly getting there getting there, and over time, we think that as we -- if you look at those layers of the stack from the hardware to the middle operating systems via applications, the value increases as we go up the stack. So we really see that shift changing towards the software as we go. You also asked about recurring, certainly that's an attractive model today both for customers as well as for software vendors, and we intend to fully participate in that model as we go forward.

Q - C.J. Muse

Very helpful. So, Omniverse, seems to have a very strong presence in many of NVIDIA's recent presentations. What is it that makes you so excited about this platform and how should we think about the opportunity of the next few years for the company I guess, either in terms of revenues, user growth or whatever metric you think is worthwhile discussing?

A - Manuvir Das

Yeah, I think the reason we are so optimistic about Omniverse CJ, is because of the reaction we've seen. So Omniverse is basically a platform for collaboration, for collaborative design. It began because as you know NVIDIA as a company does a lot of design work, 3D work itself and we have remote teams and especially with COVID-19, it was necessary to work in a remote fashion. So we saw that opportunity and we built a software platform and as we built it and started working with early customers, we saw really quite high interest.

So for example, already today, we're just launching the platform now and we have more than 400 companies already evaluating the platform with us and companies like BMW, Erickson, Bentley, some of the some of the best known companies in the world are very excited and are working with us but just like our own engineers were excited with this platform, right? So that's one reason we see this because we realize now that the problem the problem we began to solve for ourselves is actually a universal problem, because more and more companies work collaboratively, and there's a lot of design work, no matter what kind of company you are.

The other reason we're very optimistic is because when you do the market analysis, what we realized is that over 20 million designers and engineers in the world today that can benefit from a platform like this where they can work collaboratively, work

under the house and each part of the team can produce different parts of the solution that can be plugged together, right and so again we have a licensing model for Omniverse. We've announced that already announced at Computex for a typical company, they can have a starting point of about \$14,000 per year for the company to get started with a few designers and so again, if you do the math of the number of designers and engineers, we think can benefit from this technology was with our licensing model we think there is really significant opportunity here.

Q - C.J. Muse

I guess, perhaps moving to your recent announcement that involving NetApp an base command software development platform. It seems like this is an important driver to create a subscription model for DGX SuperPOD and I guess, could you maybe highlight how we should think about the opportunity to NVIDIA. So, what makes this offering unique versus something like using GPU as a service through cloud players.

A - Manuvir Das

Yes, I'm glad you brought this up CJ, because this is a very important -- a strategically a very important initiative for NVIDIA and I'm glad to do lay -- The first thing to understand is we have very strong partnerships with the cloud vendors vendors. We are very happy to work with them. They have done great work in making GPU based technology available to a wide base of customers. That's a thriving model, it will continue to thrive and we think for many, many customers that is the place where they will do the AI work and we're very happy to support all the CSPs that we can ask.

That said, we definitely see a need from our customer base to complement that with solutions that are more hybrid or private cloud and the reason for that is because it's all about data gravity. AI is based on taking the data that you have and putting it to use and for every customer, the requirements about where they data needs to reside are quite different, and data is very difficult to move. So the AI computation capability needs to be where the customers data is and there's a large section of customers whose data is not in the public cloud will not be in the public cloud, cannot be the public cloud. So we're just covering that opportunity on behalf of our customer base. So that's first thing to understand.

Second thing is there's actually two things going on here. The first is, the AI has been certainly from the outside difficult technology to operationalize, and we ourselves are in NVIDIA are one of the largest consumers of AI on the planet, because of all the work we do on self-driving cars and robotics and such and so over the years, we built for ourselves internally our software platform called Base Command that all of our data scientists use every day to do the AI work, and to utilize the equipment as well as we possibly can and every time we work with the customer to talk to them about SuperPOD or provide them equipment, we give them a little tour, and we show them a glimpse of what we've got inside the in-house with base command and the question we always get from customers is why can I not get that and it just took us a little while to productize that and get it ready for customers and we did that and so we made that available now, so that's the first thing. Base command platform is

this software that we use internally for years that we're now making available to customers and it can actually operate on equipment, like, DGX SuperPOD that we provide, or it can even operate with infrastructure in the public cloud and we announced some partnerships with public cloud providers.

The second part of this CJ is, where is the hardware that base command platform is going to help you operate and so this is where we have really taken a small step, but a big move for NVIDIA where we have said, okay, for the first time, we will take the equipment that we've got the SuperPOD, the flagship product, the instrument of AI and we will provide it to you in a subscription model where NVIDIA takes on the risk of owning the equipment. NVIDIA operate the equipment on your behalf, so you can just consume it and you can consume it for only months at a time if that's what you like to do and so that's a new endeavor for NVIDIA. We're very excited about it because we think it really lowers the barrier to entry for customers. Many more customers will experience SuperPOD, and either move on to procure their own or scale up in the cloud and how is this different from what people can do in the cloud today. The DGX SuperPOD is a unique capability. It's an instrument of AI, it's a combination of hardware, software, tuning packages that really allow you to do best-of-breed AI with all the right network and everything built-in, it's a unique capability and so it's been provided now to customers in this fashion so they can experience it for themselves.

Q - C.J. Muse

Interesting. So, just to be clear. So you're going to operate effectively this mini data center for your enterprise customers, but it will be on site at the each enterprise player.

A - Manuvir Das

Okay. So, it's a hosted offering C.J. So we certainly so you referenced NetApp we have a partnership with NetApp because they are the storage provider with infrastructure, and we're also working with Equinix, which as you well know is a very, very large provider of co-location space and services across the globe. In fact, a lot of the CSPs themselves are hosted with an Equinix. So all of these environments that we're providing hosted within Equinix data centers. They're not in the company's individual data centre.

Q - C.J. Muse

Got you. Well, I guess one last question on this front, I guess how do you think about workloads in the cloud versus on-prem for a more or through this kind of new partnership that you have with NetApp and Equinix, looking out five years and as part of that, just to respond to the question that I received, I think there's a concern out there that AI as a service will be hosted by cloud players and that will kind of minimize the amount of GPU's or other accelerators one might need. You know how would you address that as well as part of this question?

A - Manuvir Das

Yeah, that's a great question that's a two-part answer C.J. So firstly you asked about how we anticipate the mix between, AI workloads being run in the cloud versus on-premises. I think it's very important to understand that when we talk about on-premises or non-cloud, a huge part of that going forward will be the Edge, where a lot of companies do their work today. Whether you're a retail store or a factory an industrial company and certainly, when you're out of the Edge, right, then you're not doing your competition in the cloud and a lot of AI will happen in the Edge, because the data is coming in right there in real time, the data has to be analyzed and processed right there and the outcomes have to happen right there, right? So certainly, the edge will be a big dominant force as to where AI is deployed.

Secondly for enterprise customers, that own data centers versus the cloud, it's a matter of data gravity. I think in many industries, healthcare being a great example of companies that have large amounts of data certainly, it's beneficial for them to hold the data in-house. The data is really the IP and so again, it's about the computation needs to live with the data and CJ, this is why you see every cloud provider also working on non-cloud solutions, right, to move their capabilities outside of the public cloud, there's a reason for that, right because data gravity is what is going to be the dominant factor going forward. So that one part of it.

The second question regarding the concern you raised that, if it's just going to be AI as a service from cloud providers, what does that mean for NVIDIA? There's a reason why we have thousands and thousands of engineers the best engineers in the planet working every day on our hardware and software, platform, we believe that we have the best differentiated full-stack platform for AI. Cloud partners who provide AI services in the cloud today use our hardware and software quite a bit in providing these services and so that's a great thing and we fully anticipate that as the AI services in the cloud continue to develop and progress and have a footprint they will be using NVIDIA hardware and software technology as they do today and we work very closely with them. We joined the engineer with them and we cater to the requirements that they get from their customers. So we're very confident that as customers adopt the public cloud more and more for AI services, NVIDIA will be front and center in that journey with the cloud providers and the customers. And then, of course, there's a complimentary world that we are working on with the likes of NetApp and Equinix and others on-premises and the Edge.

Q - C.J. Muse

Very helpful. Perhaps we can move to the hardware side of the house, we were few quarters into a 100 introduction into the enterprise world. Can you speak to reception and anything that jumps out at you that is noteworthy and worthy of discussion?

A - Manuvir Das

Yeah, I think demand has been very strong CJ. The first thing is all of the OEMs built so and qualified systems with A100, they moved very fast as you know server companies go through a qualification cycle, there's some physics involved in that, but they saw the opportunity with us, and they really turned on a dime, and we are very thankful to them for how quickly they reacted and you can see that effect now in

the market. And then also our CSP partners adopted A100 very quickly, and the uptake has been amazing. In fact, we see very, very strong demand in the public clouds for these A100 instances that the cloud providers have. So I think that's great. In terms of the areas, we've seen very strong adoption in consumer internet, financial services, higher education. So overall, we are very pleased CJ and the uptake has been more than we had expected at the outset. So very pleased with that. Yes.

Q - C.J. Muse

When you think about the intelligent Edge NVIDIA's opportunity with the DGX platform can you speak to the work of the company has done to foster adoption here and how specifically does this ease the move of AI to the Edge?

A - Manuvir Das

Yeah, I think C.J it's a really good point to make here is that AI is really an ecosystem effort because you can't just build a chip and say, okay, there's a great chip, everybody do AI, because there's million use cases. So, you need an ecosystem of ISVs who produce the right application, right and this is especially true on the Edge because every company is trying to do something completely different on the Edge. Maybe, I'm a retail store and I'm looking at people picking up stuff from counters, and oil and gas facility and -- perimeter security. The use cases are just completely different right on a factory and in robotics and many, many years ago, at Nvidia, we decided to take a platform approach. We put programming model on top of our GPU which we call CUDA, and that's why today we have 2.5 million developers in our ecosystem who build on CUDA.

All of that philosophy that we chose as opposed to saying has built a dedicated chip for every particular use case like some kind of a PGA, kind of approach, that platform approach is really paying off when we come to the intelligent Edge, because one platform from NVIDIA, every ISV that we work within the ecosystem can quickly develop the application that is tailor-made for the particular use case that they're one. Maybe I'm working with Walmart, on a particular suite -- solution for the retail stores, or I'm working with BMW on the thing they're putting it on the factory floor, but it's the same chips, it's the same SDK, it's the same frameworks and libraries and I can rely on all of that and just add my IP and top that to provide a solution to the customer, right?

So we think this is a really strong winning formula and this is been our approach for the intelligent edge. So, we have a EGX is a hardware platform every OEM is producing ruggedized or purpose-built form factors that can go into the Edge into a closet somewhere into place where it's hard, when it's cold and then we've got the whole software stack -- that ISVs can depend on and then the final piece of the puzzle CJ, is that we built a platform called fleet command. So just as we talked about base command earlier, which data scientists can use to develop the AI, fleet command is a platform that can be used to them deploy the AI out the Edge, because if you think about the retailer for example, we work with Walmart, right and you can think about how many stores Walmart has around the world and if you're putting in a little bit of this equipment on the edge in every store in every store this is not traditional IT, where you can have IT people sitting in every store, walking up and

down watching the equipment, because there's only two or three servers in every store.

So what you really need is, you need some kind of command module that is centralized that manages all these systems in every store that keeps all the AI up-to-date that make sure nothing's been tampered with, it's all secure, and this is what fleet command is, right? So I would say the short summary of it is, we think NVIDIA has a massive opportunity with the intelligent edge, because we bet on the right platform, the universal platform, we built the ecosystem around it, and we put the tooling around with fleet command, to really help people operationalize.

Q - C.J. Muse

Very helpful. So we have about five minutes left right. I got about seven more pages of questions, so I'll try to put a couple together. So you have kind of more product you can show with the Mellanox acquisition. So you also have three chip strategy bringing in Arm-based CPU. So, I guess perhaps if you look out over the next three to five years, what gets you excited in terms of the other hardware capabilities that you are able to bring to bear and I guess having a full pre-chip solution, how will that really move the needle to the enterprise?

A - Manuvir Das

Yeah, I think it's a great way to wrap this up C.J. Really the two fronts, one is with the data processing unit of the DPU, which is the technology that Mellanox brought into Nvidia, and then the other is the CPU that we announced the -- CPU based on Arm. They're actually in some ways diametrically opposite and it's important to understand. The reason we are so excited about the DPU at NVIDIA is because we believe that going forward, every server that is shipped will have one of these DPUs in it. It is a universal opportunity. So any kind of opportunity mathematics you would do, you just take the number of servers that have been shipped and then you just put DPU in every server and the reason for that is because more and more security is -- stakes in any data center and security is becoming harder and harder to do because of the sheer volume of traffic, the amount of data and the DPU is what actually enables a security solution that can actually stand the test of time in every data center. So we believe there will be a DPU in every server.

The other reason is because today servers are getting bogged down because the CPU is being used to do a lot of the sort of the data center processing that surrounds the application and that's taking away oxygen from the applications themselves and so you need more servers to run the application than you actually should, if you had access to the full CPU and what the DPU does is it offloads that work from the CPU, gives it another place to run, so that the CPU is fully free to run the applications, right. So we see this model going forward when every server you'll have a CPU, you'll have GPUs to run the applications that are suitable for GPUs, but you'll always have a DPU as an essential component and so we're very, very excited about this opportunity CJ and we think all the work we're doing, we've been laying the groundwork with the OEMs, with VMware, with cloud providers to trigger this movement to have a DPU in every server.

Now the CPU work is in very much in contrast for this. The reason we are working on Grace is because for a specific situation, which is AI down at extreme high scale when the models are very, very large. So for example, we are now getting to models in natural language processing that are not just tens of billions of parameters, but soon we'll have trillion parameters, this is very large models and when you try to process these models, no matter how many servers you have, the communication between the CPU and the GPU becomes the big bottleneck, because you've got this engine called GPU, that can process this very large-scale AI, but it all has to come from the CPU and so we designed Grace very carefully to be a CPU that is very tightly integrated with the GPU there is alongside on the system and I'm saying this to help convey that this is a niche solution for this particular use case of high scale AI.

We believe that for the foreseeable future, the majority of enterprise workloads will be running on x86 CPUs and we work very closely with the very large system of x86 based server systems and we anticipate doing that for a long time, and I think that is going to be the dominant part of our business for time to come, but we've developed Grace to really expand the frontier, right and move the needle as to what best-of-breed -- looks like. So, the DPU is every server big, small, regardless of workload, you got to DPU -- Grace is exactly the opposite. For some small number of workloads, it's going to be the difference between being able to do it or not, and for everything else, x86 CPUs would be awesome.

Q - C.J. Muse

Great. One last question, I'll sneak again. So, the focus at the beginning was democratizing AI. We learned in the call this software opportunity is going to be at least as large as the hardware opportunity and I had a presentation earlier today with another company talked about over the next three to five years, at least the third of workloads done in compute using AI and so that would seem to create a strong path of growth here. You got 30 seconds, how do you want to close on that kind of framework?

A - Manuvir Das

Yeah. I think, I'll close by saying, we are already in a very good point, because when we think about our data center business today, it is already the case that more than half of our data center business is this enterprise business of companies within their own buildings, deploying AI technology in spite of it being a nascent area and all the complexities that come about in adopting new technology and I think we're at the beautiful point now, where we have learned so much from all from all our customers have adopted AI, two things are clear to us. One every company will benefit from AI and two, we have the platform that will make it possible and so we are just the right moment to start this journey of democratizing AI, and we working with all of the ecosystem C.J to make this happen, and that's why we're so happy about those partnerships, because this is not something that is done alone it is done in conjunction with trusted vendors that enterprise customers already work with, this is about coming to enterprise customers with solution that is easy for them to use, that is comfortable for them while providing new capability.

Q - C.J. Muse

Excellent. Well, thank you for your time, greatly appreciated, and I wish you the best of luck.

A - Jensen Huang

Thank you very much for having me and so having NVIDIA.

Q - C.J. Muse

Be well.

This transcript may not be 100 percent accurate and may contain misspellings and other inaccuracies. This transcript is provided "as is", without express or implied warranties of any kind. Bloomberg retains all rights to this transcript and provides it solely for your personal, non-commercial use. Bloomberg, its suppliers and third-party agents shall have no liability for errors in this transcript or for lost profits, losses, or direct, indirect, incidental, consequential, special or punitive damages in connection with the furnishing, performance or use of such transcript. Neither the information nor any opinion expressed in this transcript constitutes a solicitation of the purchase or sale of securities or commodities. Any opinion expressed in the transcript does not necessarily reflect the views of Bloomberg LP. © COPYRIGHT 2024, BLOOMBERG LP. All rights reserved. Any reproduction, redistribution or retransmission is expressly prohibited.