Company Participants

- Colette Kress, Chief Financial Officer
- Eyal Waldman, Chief Executive Officer
- Jensen Huang, President and Chief Executive Officer
- Simona Jankowski, Vice President of Investor Relations

Other Participants

- Aaron Rakers, Analyst, Wells Fargo.
- C.J. Muse, Evercore
- Mark Lipacis, Jefferies
- Pradeep Ramani, Analyst, UBS
- Stacy Rasgon, Bernstein Research
- Unidentified Participant
- Will Stein, Analyst, SunTrust

Presentation

Operator

Good morning. My name is Misti, and I will be your conference operator today. Welcome to the NVIDIA to Acquire Mellanox Conference Call. At this time all participants are in a listen-only mode. Later we will conduct a question-and-answer session and instructions will follow at this time. (Operator Instructions)

I would now like to turn the call over to Simona Jankowski, Vice President of Investor Relations at NVIDIA. Please go ahead.

Simona Jankowski {BIO 7131672 <GO>}

Thank you. Good morning, everyone, and good afternoon to those in Israel. With me on the call today from NVIDIA are Jensen Huang, President and Chief Executive Officer; Colette Kress, Executive Vice President and Chief Financial Officer. We're also joined by Eyal Waldman, President and Chief Executive Officer of Mellanox.

I'd like to remind you that our call is being webcast live on NVIDIA's Investor Relations website. The webcast will be available for replay following this call. In addition, you can find a presentation on the proposed transaction, as well as FAQs on the NVIDIA Investor Relations website. The content of today's call is NVIDIA's property. It can't be reproduced or transcribed without our prior written consent.

During this call, we may make forward-looking statements based on current expectations. These are subject to a number of significant risks and uncertainties, and our actual results may differ materially. For a discussion of factors that could affect our results and business, please refer to our most recent Forms 10-K and 10-Q, and the reports that we may file on Form 8-K with the Securities and Exchange Commission. All our statements are made as of today March 11, 2019; based on information currently available to us. Except as required by law, we assume no obligation to update any such statements.

With that let me turn the call over to Jensen.

Jensen Huang {BIO 1782546 <GO>}

Thanks, Simona. And thanks everyone for joining on short notice. Today, we're announcing our proposed acquisition of Mellanox for $6.9 billion. Mellanox is headquartered in Israel and is world-renowned for its high performance datacenter interconnect technology. We have worked together for many years and have partnered to build many of the world's highest performance datacenter systems.

We are combining two of the world's leading datacenter and high performance computing technology companies today. Between us, we power over 250 of the world's top 500 supercomputers, and we name every major computer maker and cloud service provider as our customers. Let me tell you why this makes sense for NVIDIA and why I'm so excited about it?

As you know, NVIDIA is focused on high-performance computing. Because of workloads such as artificial intelligence, scientific computing and data analytics, performance demands on hyperscale and enterprise datacenters are skyrocketing. From a computer architecture point of view, two dynamics are particularly important. First, CPU performance advancement is slowing just as computing demand is skyrocketing. As a result, accelerated computing and accelerated networking, or otherwise known as CPU offloading, is a fundamental path forward.

Second, as data and compute-intensive workloads like AI and data analytics are growing exponentially, they require thousand -- tends of thousands of server nodes working together. To support these new workloads, the computing nodes and the connecting fabric will form a datacenter-scale compute engine. Optimal design will translate directly to higher datacenter throughput and utilization and lower operating cost.

To accelerate applications, NVIDIA optimizes across the entire application stack from architecture to chips, to systems to algorithms. In the future, we want to optimize datacenter-scale workloads, again, across the entire stack from the compute node to networking to storage. For this reason, Mellanox's system-to-system datacenter-scale interconnect technology is important to us. We believe our platform would be stronger and deliver the best possible performance for datacenter customers by innovating across the computing, networking and storage stacks.

Beyond these significant product synergies, we're excited about the tremendous cultural fit with Mellanox. Like NVIDIA, they have performance-driven culture with a focus on innovation. I look forward to working even more closely with their talented employees and visionary leaders to invent the future of computing. Great to have you on the call Eyal.

Eyal Waldman {BIO 3215622 <GO>}

Thank you, Jensen. I am delighted to the vision that Jensen has set out for accelerated computing is one that we fully share at Mellanox. The future of computing will benefit from holistic innovation across the entire technology stack from compute, to networking, to storage. Combining our two companies is a natural extension of our existing partnership and a great feat, given our common performance-driven cultures.

We have built Mellanox-based on technical excellence, innovation and staying nimble even as we've grown, and NVIDIA is an excellent match for us. Combining our companies will be good for our customers, who will benefit from accelerated innovation in high-performance computing, good for our partners, who will benefit from optimized solutions across the entire computing and networking stack, and good for our employees, who will have new growth opportunities and challenges.

Let me now turn it over to Colette to give more to give more details on the proposed transaction. Colette?

Colette Kress {BIO 18297352 <GO>}

Thanks, Eyal. Today we are announcing the acquisition of Mellanox for $125 per share in cash, implying a total enterprise value of $6.9 billion. We expect the acquisition to close by the end of calendar 2019, subject to regulatory approvals in various jurisdictions, as well as approval by Mellanox's shareholders. The transaction has been approved by both NVIDIA's and Mellanox's Board of Directors.

In addition to the strategic rationale Jensen outlined, we believe this is also a financially compelling transaction. The deal is expected to be immediately accretive to non-GAAP gross margins, non-GAAP earnings per share, and free cash flow, even before assuming any cost or revenue synergies.

For those less familiar with the company, Mellanox generated approximately $1.1 billion in revenue in calendar 2018, up 26% from the prior year. Calendar 2018 GAAP gross margin was 64.3%, and non-GAAP gross margin was 69.2%. GAAP operating income was $112 million and non-GAAP operating income was $270 million.

Mellanox's customer base and channel partners have significant overlap with our own, and customers include large hyperscale and cloud service providers, as well as many of the world's largest super-computing centers, direct customers, including many of the same large IT infrastructure OEMs that we partner with for our go-to-market. We believe that there will be opportunities to cross-sell and enhance our

broader solution set over time given each company's technology assets and customer bases in high-performance computing, enterprise, and hyperscale. But we are conservatively assuming no initial revenue synergies.

We are also not modeling any cost synergies at this time. It is a very strategic deal and we highly value Mellanox's employees. We will continue to invest to go after the large (inaudible) we see for the combined company.

Lastly, let me touch on capital allocation. We view the proposed acquisition to be a terrific use of our cash, not because it's immediately accretive from a financial perspective, but most importantly, because it is a strong strategic fit from a technology and cultural perspective. Further, this announcement does not change our capital return program for the year. Of the $3 billion we intend to return to shareholders by the end of fiscal 2020, we've returned $700 million through share repurchases during the fourth quarter of fiscal 2019. We still plan to return the remaining $2.3 billion to shareholders through the end of fiscal 2020 in the form of dividends and buybacks.

We will now open the call for questions. Operator, would you please poll for questions? Thank you.

Questions And Answers

Operator

(Operator Instruction) And your first question is from C.J. Muse with Evercore.

Q - C.J. Muse

Question -- I guess Jensen, first question, can you talk a bit about the strategic rationale here, particularly around which end markets you're looking to target? HPC is obviously the biggest potential, but would also love to hear is public cloud the main focus here, is it more of offering a platform solution now like DGX into the enterprise?

A - Jensen Huang {BIO 1782546 <GO>}

Yes. C.J., That's a great question. The datacenter market today consists of high performance computing or supercomputing, which is the initial target of our companies CUDA architecture. It also includes hyperscale for internal consumption of Internet services, provisioning of Internet services, cloud service providers offering them as cloud infrastructure and then enterprise computing. These are really the major segments if you will from the way we look at datacenters.

As you know that one of the most important drivers in the datacenter world today is a new type of workload that's emerged and it's a data-centric type of workload. It requires an enormous amount of data and a huge amount of computation to infer insight, to extract insight out of data. We call it artificial intelligence and machine

learning and data analytics. This type of workload is just simply too big to fit on any one computer. And there's a couple of ways that you could solve the problem, but the dynamics are basically related to, on the first hand, CPU scaling has slowed. And so, without offloading the algorithm somehow and accelerating it using new accelerators like what we do with GPU computing and what Mellanox does with accelerated networking, it's really not possible to continue to scale computing as Moore's Law ends.

The second is, because the data size and the compute size is so large, it won't fit in any one computer, and so, thousands of computers have to be networked together to work together. The dynamic that is happening here is that, in the future, it won't just be server-scale computing that people do, but it will be datacenter-scale computing, where the network becomes an extension of the computing fabric. You see early trends of that already, software defined networks, intelligent networks, but long-term, even some part of the computation will go into the network. And so, where the compute node starts and ends and where the networking starts and ends, as part of the overall computing fabric will be very difficult to discern and we'd like to be able to -- we want to take a lead in architecting and developing solutions that innovate across all of these stacks.

Now, of course, the solutions will have to be open because everybody's datacenter -- everybody's datacenters are different and everybody's amount of supporting traditional software stacks or new software stacks like containers and Kubernetes are going to be a little bit different and datacenters will be heterogeneous for a very long time because of that. And so our ability with Mellanox to be able to support supercomputing networks from InfiniBand, all the way to hyperscale datacenter networks with Ethernet, and to be able to put acceleration and intelligence into the entire fabric is really, really powerful. And so I hope that answers your question. But basically, we would like to accelerate growth into our datacenters by being able to architect across both the compute as well as the network fabric.

Q - C.J. Muse

That's very helpful. And if I could ask a quick follow-up, Colette, you talked about no cost synergies at this time. Are there any cost savings from the combination of two public companies, et cetera, that we should be thinking about?

A - Colette Kress {BIO 18297352 <GO>}

Yes. Thanks for the follow-up question, C.J. As we discussed, this acquisition is certainly about driving accelerated growth and improving our position in the datacenter. Mellanox is well managed and it already has quite a lean organization. We have many areas of operational leverage that we can accelerate for both of our development plans as we move forward. And over the long term, we certainly see opportunities for operating leverage of bringing these two companies together.

Q - C.J. Muse

Thank you.

Operator

And your next question is from Aaron Rakers with Wells Fargo.

Q - Aaron Rakers {BIO 6649630 <GO>}

Congratulations on the transaction. One housekeeping item and then one architectural question. First of all, Colette, just any color on, is there a breakup fee, the size of the breakup fee, any details on that? And then on the architectural question, there is a lot of discussion around the role of InfiniBand in the context of HPC and I understand what you just said around bringing intelligence into the networking. I'm interested, Jensen, how you see Mellanox's Ethernet switch business particularly playing into the portfolio, given the traction that the company has started to see over the last couple of quarters? Thank you.

A - Colette Kress {BIO 18297352 <GO>}

So, let me start first. We're going to concentrate over the rest of the calendar year to focus on closing this deal. We need regulatory approval in the US and China and a couple of other jurisdictions, but this time, we think that's our right focus before closing.

A - Jensen Huang {BIO 1782546 <GO>}

Yeah. So, first of all, the easiest way to think about that, is that, I think the world of their Ethernet solutions, and the reason for that and the reason why we -- as you know, we work with them in high-performance computing datacenters all of the world, as -- CSPs all over the world. We have our product called T4, which is about scale-out accelerated computing. And wherever we do scale-out accelerated computing, scale-out meaning, small number of GPUs accelerators inside a large number -- inside each node and many nodes are being connected together, build a little bit differently than DGX for example. DGX is a scale-up architecture. And so T4 was designed to be a universal accelerator for hyperscale datacenters and it's intended for scale-out. In those kind of environments, we see -- we work with Mellanox a lot. And the reason for that is because of the software stack.

It turns out, that the best way to accelerate a datacenter is to offload the CPU, to keep us little of the data movements from being copied around as possible. And so the software stack, the RDMA software stack and the RoCE software stack that Mellanox has created over the years, and the body of work that they've developed in accelerating the networking stack, is just world-class. That's where you get the greatest value of extending the compute node and the networking nodes -- networking fabric to become an extended computing fabric. You want to make sure that as much of the software is optimized as possible. And they do this really fantastically throughout.

Now, we're going to continue to work with switch companies all over the world and the reason for that is, of course, datacenter is heterogeneous by nature, there are parts of the datacenter where certain partners are going to be particularly great at. Arista is great in the segments that they serve and Cisco continues to be incredibly

valuable. And so we have to work together as an industry, and our goal is to make sure that the computing fabric is as optimized as possible and Mellanox plays a really important role in that. And so, I love all of their product lines, and particularly, I love their software stack and their architecture, on which the software stack runs on is just extraordinary.

Q - Aaron Rakers {BIO 6649630 <GO>}

Very good. Thank you.

Operator

And your next question is from Timothy Arcuri from UBS.

Q - Pradeep Ramani {BIO 19683324 <GO>}

Hi, gents. This is Pradeep Ramani calling for Timothy Arcuri. I think you just touched on the software stack aspect of it and that's exactly what we wanted more details on. So from a software perspective. I guess, how much more work do you have to do in terms of closely integrating RDMA and RoCE with NVIDIA products? And can you talk about the capabilities that you have in-house in terms of the software engineering capabilities to do that and just the roadmap in general?

A - Jensen Huang {BIO 1782546 <GO>}

Yes. As you know, we do a lot of software inside the Company even for networking and the reason why we do it is because we do it for NVLink, which is a processor-to-processor in system interconnect. And Mellanox develops system-to-system interconnect. NVLink and Mellanox are complementary and the reason why we work so closely together is because of the complementary nature of these two interconnects. And so, we have a great deal of expertise in both companies in how to manage networking and how to optimize the networking stack for the highest throughput and the ability to connect the largest number of computers. And so, both companies has a great deal of expertise in this area.

The important thing is, at the highest level, these are two -- we are two of the world's leading high-performance computing technology companies. Between the two of us, we power over 250 of the world's top 500 computers. And so, there's a great deal of expertise in both companies. I would love to have Eyal answer some of that question too just so that we can have him do some talking, Eyal is a visionary leader and he has, as you guys know, built Mellanox from the ground up and I'll love it, if he would take part of it.

A - Eyal Waldman {BIO 3215622 <GO>}

Thank you, Jensen for the warm words. Yes. So I think there's a lot of synergy in the architecture between the GPUs that NVIDIA designs and sells and the interconnect that Mellanox designs and sells. We can now move onto the next stage of integration and build more efficient datacenters for wide number of applications and markets. The way we've done it now is through external connectivity. The next stage of building smarter, more efficient and more productive datacenters will be higher

integration of mechanisms between the interconnect and CPU, GPU that now we'll be able to do more intimately.

Q - Pradeep Ramani {BIO 19683324 <GO>}

Okay. Thank you.

Operator

And your next question is from Stacy Rasgon with Bernstein Research.

Q - Stacy Rasgon {BIO 16423886 <GO>}

Hi, guys. Thanks for taking my questions. First, I was curious how long you've been exploring this purchase in earnest before the announcement today, was it weeks or months, like how long have you actually been in this process before actually closing on the deal today?

A - Jensen Huang {BIO 1782546 <GO>}

Yeah. That's great questions. Eyal knows I've loved Mellanox for years. I have loved Mellanox for years for several reasons. One, we collaborate. We've been working together for very long time and we've built high performance computers together. They are in every one of the DGX servers, the AI supercomputers that we make. So, in our own products they are in every single one. They're all of our datacenter, the ones that we use to achieve and deliver the world's highest performances on MLPerf. We work hand-in-hand with their engineers to achieve those results all the time. We've been working together for years. And so, I've loved this company for a very long time.

I believe for a very long time also that NVIDIA will go from me a chip-level company and -- to a server level company to essentially datacenter-scale company, and then when we think about architecture even from that perspective. And so, I've loved the company for a long time. And in terms of when the acquisition talks first happened, that's going to be in the proxy and -- but obviously, it was more recent.

Q - Stacy Rasgon {BIO 16423886 <GO>}

Okay. Just to follow-up, maybe housekeeping. How are you going to account for Mellanox's revenues when you report, are you going to leave them, split out in their own segment, or are you going to put them into your datacenter segment?

A - Colette Kress {BIO 18297352 <GO>}

Hey, Stacy. This is Colette. I'll try and answer that question. At this time, we're not changing our guidance that we just provided about a month ago for the current fiscal year. When this deal actually closes, I will be able to give you more color regarding what we see going forward in terms of guidance with the two companies together and we'll talk about the overall reporting at that time.

Q - Stacy Rasgon {BIO 16423886 <GO>}

Okay. Thank you, guys.

A - Jensen Huang {BIO 1782546 <GO>}

Thanks, Stacy.

Operator

Your next question is from Will Stein with SunTrust.

Q - Unidentified Participant

Thanks for taking my question and congratulations on this transaction. Jensen, I'm hoping you can put this acquisition in sort of a historical perspective for us relative to the other ones you've done. What lessons would you say you learned from the prior acquisitions that's sort of encouraged this transaction and what lessons might you have learned from the prior deals that might provide a different approach or different tactics around this one? Thank you.

A - Jensen Huang {BIO 1782546 <GO>}

Yeah. From a scale perspective -- from a proportional -- from a scale perspective, this is surely the largest one. From a percentage or proportional size perspective, our integration, our acquisition of 3dfx was the largest and -- at the time. And at the core of the core, it's really about several things. One, it has to be -- and just as a statement, almost every single 3dfx engineer and employee that has joined us at that time are still here and many of them are on NVIDIA's management team today on the staff [ph] today. And we wouldn't be here if not for the incredible work that the people did that came to us. We are just such a different Company as a result of that acquisition.

And I think this acquisition would also be transformative for us, changing us from a from a chip and a system-level company, to a datacenter-scale level company. And several things that I learned. One, there has to be an alignment about the purpose of companies, the culture of the company, and just the core attitude, the principles that govern the companies. And I've worked with Eyal for a long time, we've worked with Mellanox for a very long time, and our two companies are driven in very similar ways. We're both best in our class. We're singularly focused on high-performance computing, we're very performance-oriented and you just can't help and love the fact that this company is also based Israel, so that the population of great technical excellence is high. And so I'm thrilled about that part of it. And so, I think number one is the alignment of purpose and culture.

The second is just alignment of vision of the future of the industry. We both see it the same way, that this transaction is about uniting the two of the leading technology companies in high-performance computing, it's about doubling down it's datacenters, it's about inventing the future of datacenters, which is going to be workloads that run across the entire datacenter. Imagine one application that's split

up in containers that's operating in servers all over the datacenter and results are being sent around from one server to another server over the network fabric. The entire datacenter becomes one giant compute engine. And this is the future. And so, where compute starts and networking starts and ends, it's going to be harder to see and it's going to be one continuous compute plain. And so the second is the alignment of the future, vision of the future of the industry.

And then, thirdly, it has to make all the economic sense that the strategy would suggest. This deal is going to be immediately accretive. You have two companies that are growing. Our strategic tenants are still fundamentally sound. We believe that the future is going to be about acceleration. Mellanox's business is growing and they're growing fabulously and for all the dynamics, the reasons that I've already outlined is the reason why they are growing. And so we have two growth companies coming together, two best-in-class companies coming together, two cultures that were built from the ground up coming together, and we see a common future, a common vision of the future.

Q - Will Stein {BIO 15106707 <GO>}

Thank you.

Operator

And your next question is from the line of Mark Lipacis from Jefferies.

Q - Mark Lipacis {BIO 2380059 <GO>}

Hi. Thanks for taking my question. Congrats on the,announcement. I had two questions, one software and one interconnect. On the software side Jensen, you've discussed transforming NVIDIA from a chip company to a platform company for a while. And I'm wondering if the Mellanox acquisition necessarily mean that you're not going to be extending what you've done with CUDA on GPUs to a higher level software capability then you had before with what will become the combination of Mellanox and NVIDIA? That's the softer question. On the interconnect question, you talked about system-to-system, is it -- does the Mellanox interconnect capability, does that extend down to the chip-to-chip level and does it -- does it mean that you consider multi-chip packaging architectures in the future, is that on the roadmap? I'll stop there. That's all I have. Thank you.

A - Jensen Huang {BIO 1782546 <GO>}

Yes. I'll go backwards. NVLink is unquestionably the best processor-to-processor interconnect that world has today. It is just incredibly high-speed and it's designed so that CUDA programs could view all of the processors that are connected together as if one large computing engine. And so, NVLink has a very special purpose and it's really about connecting processors to processors within one system. Mellanox is about connecting multiple systems also with the same sensibility of very low latency, very high throughput and just incredibly well optimized software stacks that are integrated with the gigantic body of software that's running on top of the datacenter. If you look at their body of work, whether it's in high-performance computing, supercomputing, enterprise computing or in hyperscale computing, their software

stack is integrated deeply into the world's body of work if you will. And so, I think. I think their interconnect and our interconnect, their software and our software, has very specific places to exist, and they're going to evolve from there.

At the higher level, it is very clear that compute and networking will be one large connected compute plain, if you will, that the networking will be an extension of the compute fabric, or another way to say it is that, in the future, computing will be done not just within the server, but computing of an application will be done at the datacenter-scale. One application will run across the entire datacenter. It's either all of the servers coming together, all of the nodes coming together, working together to solve one problem or the application has been broken apart into a whole bunch of little containers with AI programs running within it, and the results of one as been essentially email to -- if you will, sent to the input at the [ph] other. And so, the traffic inside datacenters has exploded because of those phenomenon. And -- but at the highest level, the easiest way to think about that is, in the future, we're looking at datacenter-scale compute engines, the servers will be connected by networks that will create a compute plain that is continuous and the best architecture is going to have to contemplate this holistic nature.

Q - Mark Lipacis {BIO 2380059 <GO>}

That's helpful. Thank you.

A - Jensen Huang {BIO 1782546 <GO>}

This is the last question? Okay. You guys, this is the last question. Thanks for joining us. This was a real -- this is a really exciting day for our Company. And I'm super delighted that we were able to acquire Mellanox today and announce the acquisition of Mellanox today, and this is going to be the beginning of a new NVIDIA that I look forward to sharing with you as time goes. Thank you.

Operator

And this now concludes today's conference call. You may now disconnect.